



STA 426 First Lecture

- survey
- course structure
- Molecular Biology lecture (Hubert)
- R survey + computing + Exercise 1



Today's structure

9.00-9.45: Survey + Course Structure (Mark)

10.00-10.45: Introduction to Molecular Biology (Hubert)

11.00-11.45: Computing + R quiz + Rmarkdown exercise





Survey: Statistical Insight

app.klicker.uzh.ch/join/marobi




Question 1: Which plot highlights more (statistical) evidence for a change in the population means (between orange and blue)?





Question 2: In your view, what best describes the associations shown in the plot of 'x' and 'y' ?




Question 3: Of these equations, which one resembles the standard two sample t-test ?

1
$$\frac{(\hat{p}_1 - \hat{p}_2)}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

2
$$\sum^k \frac{(\text{observed} - \text{expected})^2}{\text{expected}}$$

3
$$\frac{(\bar{x}_1 - \bar{x}_2) - d_0}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}},$$

Question 4: Given these boxplots, which of two underlying distributions are more similar?







Question 6: Given this design matrix, describe the experimental design.

$$X = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$

Question 8: Which plot highlights more (statistical) evidence for a change in the population means (between orange and blue)?





Course communication

- Video situation: we will have recordings, hopefully within 24 hours of the lecture
- Slack: vast majority of communication happens here (note: all invitations were sent to UZH email addresses)
- Except for exceptional circumstances, **no emails please**; communicate on Slack only (maybe later on GitHub)
- Slack policy: unless private, ask questions in a public channel (please note: *good questions get good answers*); use threads when relevant; good manners/behaviours are expected



Course evaluation

1. Journal club presentation	20%
2. Project	50%
3. Exercises	30%
4. Technology day (participation)	0% or -10%



Rough structure of lecture/exercise time

Monday mornings: we will run X.00-X.45; X in {9,10,11}

- Lectures and Exercises
- Lecture/journal club presentation (9.00-whenever)
- Remaining time: free (can be used to work on exercises; we are available for questions)



M.Sc. thesis projects

If you are:

- in a M.Sc. programme (ETHZ or UZH)
- have a solid background/experience in mathematics / statistics / computation
- have an interest in research in this field (“statistical bioinformatics”)
- looking for a thesis project

→ Discuss a project in my lab

Critical skills needed by statisticians (Jeffrey Leek's words):

With all the excitement going on around statistics, there is also increasing diversity. It is increasingly hard to define “statistician” since the definition ranges from [very mathematical](#) to [very applied](#). An obvious question is: what are the most critical skills needed by statisticians?

So just for fun, I made up my list of the top 5 most critical skills for a statistician by my own definition. They are by necessity very general (I only gave myself 5).

1. **The ability to manipulate/organize/work with data on computers** - whether it is with excel, R, SAS, or Stata, to be a statistician you have to be able to work with data.
2. **A knowledge of exploratory data analysis** - how to make plots, how to discover patterns with visualizations, how to explore assumptions
3. **Scientific/contextual knowledge** - at least enough to be able to abstract and formulate problems. This is what separates statisticians from mathematicians.
4. **Skills to distinguish true from false patterns** - whether with p-values, posterior probabilities, meaningful summary statistics, cross-validation or any other means.
5. **The ability to communicate results to people without math skills** - a key component of being a statistician is knowing how to explain math/plots/analyses.



Learning outcomes (in my words)

- Understand the fundamental “scientific process” in the field of Statistical Bioinformatics
- Be equipped with the skills / tools to preprocess genomic data (Unix, Bioconductor, mapping, etc.) and ensure reproducible research (R / markdown)
- Have a general knowledge of (some) **types** of data and **biological applications** encountered with high throughput genomic data
- Have the general knowledge of the range of statistical methods that get used with microarray and sequencing data
- Gain the ability to apply statistical methods / knowledge / software to a collaborative biological project
- Gain the ability to critical assess the statistical bioinformatics literature
- Write a coherent summary of a bioinformatics problem and it’s solution in statistical terms



The semester-long course structure (subject to change)

Date	Lecturer	Topic	Exercise	JC1	JC2
19.09.2022	Mark + Hubert	admin; mol. bio. basics	R markdown; git(hub)		
26.09.2022	Mark	interactive technology/statistics session	group exercise: technology pull request		
03.10.2022	Hubert	NGS intro; exploratory data analysis	EDA in R		
10.10.2022	Mark	limma + friends	linear model simulation + design matrices		
17.10.2022	Hubert	mapping	Rsubread		
24.10.2022	Hubert	RNA-seq quantification	RSEM	X	X
31.10.2022	Mark	edgeR+friends 1	basic edgeR/voom	X	X
07.11.2022	Mark	edgeR+friends 2	advanced edgeR/voom	X	X
14.11.2022	YYY	hands-on session #1: RNA-seq	FASTQC/Salmon/etc.	X	X
21.11.2022	Hubert	single-cell 1: preprocessing, dim. reduction, clustering	clustering	X	X
28.11.2022	YYY	hands-on session #2: cytometry	cytof null comparison	X	X
05.12.2022	Mark	single-cell 2: clustering, marker gene DE	marker gene DE	X	X
12.12.2022	YYY	hands-on session #3: single-cell RNA-seq (cell type definition, differential state)	full scRNA-seq pipeline	X	X
19.12.2022	Mark	loose ends: HMM, EM, robustness	segmentation, peak finding	X	X



Expectations: **journal club** presentation

- 20-25 minutes (+5 minutes discussion)
- MUST:
 - ➔ be a paper about a **statistical** method in bioinformatics
 - ➔ be approved by Mark/Hubert
- Should:
 - ➔ describe the biological context and/or data collected
 - ➔ describe the (new) model used
 - ➔ describe comparisons to existing methods
- Should not:
 - ➔ be one of the papers discussed in detail in lectures: limma, edgeR, DEXSeq, etc.
- (since 2017) feedback forms from fellow students



Expectations: **project**

- ~10-15 page report, with R code in line (e.g. **quarto**)
- Describe the biological setting, statistical analysis, exploratory analysis with publication-quality graphics embedded
- Three possibilities:
 - Comparison of statistical methods (simulation / reference data + metrics)
 - Reproduce an analysis from a paper from the raw data
 - Real collaborative project with FGCZ or a local laboratory
- Be strategic: work on something related to your interests!
- Typically due at end of first working week of January



Expectations: **exercises**

- There will be an exercise **every** week
- Across 14 weeks, the *best 9* exercises are counted towards the 30%



Soft technical skills needed (developed) in this course ...

- **Data Science!**
- Use unix-like operating system to run command-line programs
- Options:
 - use your own computer (if Windows, use cygwin)
 - use renkulab.io
- R: from the command line or RStudio (<https://rstudio.com/>); getting help; creating workflows; how to make publication-quality graphics (ggplot2); knitr/Rmarkdown
- Bioconductor – www.bioconductor.org
- git/github
- bioconda/Docker (cloud computing)



Hubert's lecture




Demos:

- [git/github](#)
- [renkulab.io](#) (fork a project, start a session)
- [quarto.org](#)




```
git clone
git pull
git status
git branch
git commit
git add
git checkout
git push
```

Quick intro to Git/Github (version control)



Branching





Exercise 1

Part a: GitHub

Part b: quarto



Note: all homework submissions occur via github

Week 1 Exercise Part A:

1. Recommend: get R 4.2.1, latest RStudio, git, quarto, etc.
2. If you haven't already, create an account at github.com/join; give GitHub username (+details about computing) to Mark via <https://forms.gle/sc7ci6jPFweBE8xKA>
3. Acquaint yourself with git / github (gitlab) [1]; make sure you can check in (push) / out (pull / clone) files from command line or app [2].
4. Create a new public github repository, add a README.md (using markdown [3]) and add some content; include an image; include a web link, etc.
5. Add an Issue to the 'material' repo [4] with a link to your repo (you can delete the repo after I've closed the issue, if you want)

[1] <https://gist.github.com/andrewpmiller/9668225>

[2] <https://confluence.atlassian.com/stash/basic-git-commands-278071958.html>

[3] <http://markdowntutorial.com/>

[4] <https://github.com/sta426hs2022/material>



Rmarkdown / knitr for executable documents / reproducibility

Week 1 Exercise Part B:

1. Test your R knowledge here: <https://forms.gle/FFHiFx8UHrBVGv2R9> (only 9 questions)
2. Acquaint yourself with quarto for executable documents [1].
3. Create an HTML document with R code that samples 100 values from a negative binomial distribution (say, $\mu=10$, dispersion=2; using the parameterisation with mean= μ and variance= $\mu+\mu^2 \cdot \text{dispersion}$); create a histogram of sampled data on both the linear and log [or maybe $\log(x+1)$ due to zeros] scale; Write 1-2 sentences to describe your steps (ideally also with section headings) and report the mean and variance of the sample *in line* in the text.
4. Add the QMD and HTML files to your repo from Week 1 Exercise Part A.

[1] <https://quarto.org/>